

# LECTURE 4

## VECTOR CLOCK SYNCHRONIZATION

Prof. D. S. Yadav  
Department of Computer Science  
IET Lucknow

# PROBLEMS WITH TOTAL ORDERING

- A linearly ordered structure of time is not always adequate for distributed systems
  - captures dependence of events
  - loses independence of events - artificially enforces an ordering for events that need not be ordered.
- Mapping partial ordered events onto a linearly ordered set of integers it is *losing information*
  - Events which may happen simultaneously may get different timestamps as if they happen in some definite order.
- A partially ordered system of *vectors* forming a *lattice* structure is a natural representation of time in a distributed system

# WHY DO WE NEED GLOBAL CLOCKS?

- For causally ordering events in a distributed system
  - **Example:**
    - Transaction T transfers Rs 10,000 from S1 to S2
    - Consider the situation when:
      - State of S1 is recorded after the deduction and state of S2 is recorded before the addition
      - State of S1 is recorded before the deduction and state of S2 is recorded after the addition
- Should not be confused with the clock-synchronization problem

## VECTOR TIME (FIDGE/MATTERN/SCHMUCK)

- The system of vector clocks was developed independently by Fidge, Mattern and Schmuck.
- In the system of vector clocks, the time domain is represented by a set of  $n$ -dimensional non-negative integer vectors.
- Each process has a clock  $C_i$  consisting of a vector of length  $n$ , where  $n$  is the total number of processes.
- $VTP_i[1..n]$ , where  $VTP_i[j]$  is the local logical clock of  $P_i$  and describes the logical time progress at process  $P_j$ .

# ORDERING OF EVENTS

Lamport's *Happened Before* relationship:

For two events  $a$  and  $b$ ,  $a \rightarrow b$  if

- $a$  and  $b$  are events in the same process and  $a$  occurred before  $b$ , or
- $a$  is a send event of a message  $m$  and  $b$  is the corresponding receive event at the destination process, or
- $a \rightarrow c$  and  $c \rightarrow b$  for some event  $c$


# CAUSALLY RELATED VERSUS CONCURRENT

## Causally related events:

- Event  $a$  causally affects event  $b$  if  $a \rightarrow b$

## Concurrent events:

- Two distinct events  $a$  and  $b$  are said to be concurrent (denoted by  $a \parallel b$ ) if  $a \not\rightarrow b$  and  $b \not\rightarrow a$




A space-time diagram

# LAMPORT'S LOGICAL CLOCK

*Each process  $i$  keeps a clock  $C_i$*

- Each event  $a$  in  $i$  is time-stamped  $C_i(a)$ , the value of  $C_i$  when  $a$  occurred
- $C_i$  is incremented by 1 for each event in  $i$
- In addition, if  $a$  is a send of message  $m$  from process  $i$  to  $j$ , then on receive of  $m$ ,  
$$C_j = \max (C_j, C_i(a)+1)$$

# HOW LAMPORT'S CLOCKS ADVANCE





## IN LAMPORT'S CLOCK

- if  $a \rightarrow b$ , then  $C(a) < C(b)$
- $\rightarrow$  is a partial order
- Total ordering possible by arbitrarily ordering concurrent events by process numbers

# LIMITATION OF LAMPORT'S CLOCK

$a \rightarrow b$  implies  $C(a) < C(b)$

**BUT**

$C(a) < C(b)$  doesn't imply  $a \rightarrow b$  !!

*So not a true clock !!*

## VECTOR TIME: THE PROTOCOL

- A process  $P_i$  ticks by incrementing its own component of its clock

$$C_i[i] = C_i[i] + 1$$

- The timestamp  $C(e)$  of an event  $e$  is the clock value after ticking.
- Each message gets a piggybacked timestamp consisting of the vector of the local clock.
  - The process gets some knowledge about the other process' time approximation.

# LOGICAL CLOCKS : VECTOR CLOCK

- ❑ Vector Clock uses a vector of Integers of size  $N$ , where  $N$  is number of processes in system.
- ❑ Process  $P_i$  maintains a vector clock  $VT_i$ .
- ❑  $VT_i[i]$  is process  $P_i$ 's own logical time.
- ❑  $VT_i[j]$  is process  $P_i$ 's best knowledge of time at process  $P_j$ .
- ❑ Proposed by Fidge and Mattern and based on Lamport's scalar clocks

## VECTOR CLOCKS : UPDATE RULES

**Each process  $P_i$  has a clock  $C_i$ , which is a vector of size  $n$**

**The clock  $C_i$  assigns a vector  $C_i(a)$  to any event  $a$  at  $P_i$**

### Update rules (Internal & Message Events)

[IR1] : Clock  $C_i$  is incremented between two successive events in process  $P_i$


$$C_i[i] = C_i[i] + d$$

[IR2] When  $P_i$  sends a message  $m$  to Process  $P_j$ , it piggybacks a logical timestamp  $t$  which equals the time of the send event :  $t(m) = C_i$

When executing a receive event at  $P_j$  where a message with timestamp  $t$  is received, the clock is advanced


- $C_j[J] = C_i[J] + d$   
 $C_j[k] = \max(C_j[k], t_m[k])$  for all  $k$

# VECTOR CLOCKS EXAMPLE




Evolution of vector time.

# APPLYING VECTOR CLOCKS TO BROADCAST SYSTEM




# APPLYING VECTOR CLOCKS TO BROADCAST SYSTEM







# APPLYING VECTOR CLOCKS TO BROADCAST SYSTEM



# APPLYING VECTOR CLOCKS TO BROADCAST SYSTEM



# APPLYING VECTOR CLOCKS TO BROADCAST SYSTEM



# PARTIAL ORDER BETWEEN TIMESTAMPS

For events a and b with vector timestamps  $t^a$  and  $t^b$ ,

- Equal:  $t^a = t^b$  iff  $\forall i, t^a[i] = t^b[i]$
- Not Equal:  $t^a \neq t^b$  iff  $\exists i, t^a[i] \neq t^b[i]$
- Less or equal:  $t^a \leq t^b$  iff  $\forall i, t^a[i] \leq t^b[i]$
- Not less or equal:  $t^a \not\leq t^b$  iff  $\exists i, t^a[i] > t^b[i]$
- Less than:  $t^a < t^b$  iff  $(t^a \leq t^b \text{ and } t^a \neq t^b)$
- Not less than:  $t^a \not< t^b$  iff  $\neg(t^a \leq t^b \text{ and } t^a \neq t^b)$
- Concurrent:  $t^a \parallel t^b$  iff  $(t^a \not< t^b \text{ and } t^b \not< t^a)$

# CAUSAL ORDERING


- $a \rightarrow b$  iff  $t_a < t_b$
- Events  $a$  and  $b$  are causally related iff  $t_a < t_b$  or  $t_b < t_a$ , else they are concurrent
- Note that this is still not a total order

# USE OF VECTOR CLOCKS IN CAUSAL ORDERING OF MESSAGES

- If  $\text{send}(m1) \rightarrow \text{send}(m2)$ , then every recipient of both message  $m1$  and  $m2$  must “deliver”  $m1$  before  $m2$ .
  - “deliver” – when the message is actually given to the application for processing

# FIFO ORDER

## ORDERING PROPERTIES(1)




If a particular process broadcasts a message  $M_1$  before it broadcasts a message  $M_2$ , then each recipient process delivers  $M_1$  before  $M_2$ .

□ Message  $M_1$  ( ..... ) shows violation of FIFO order.

# Local Order

## Ordering Properties(2)



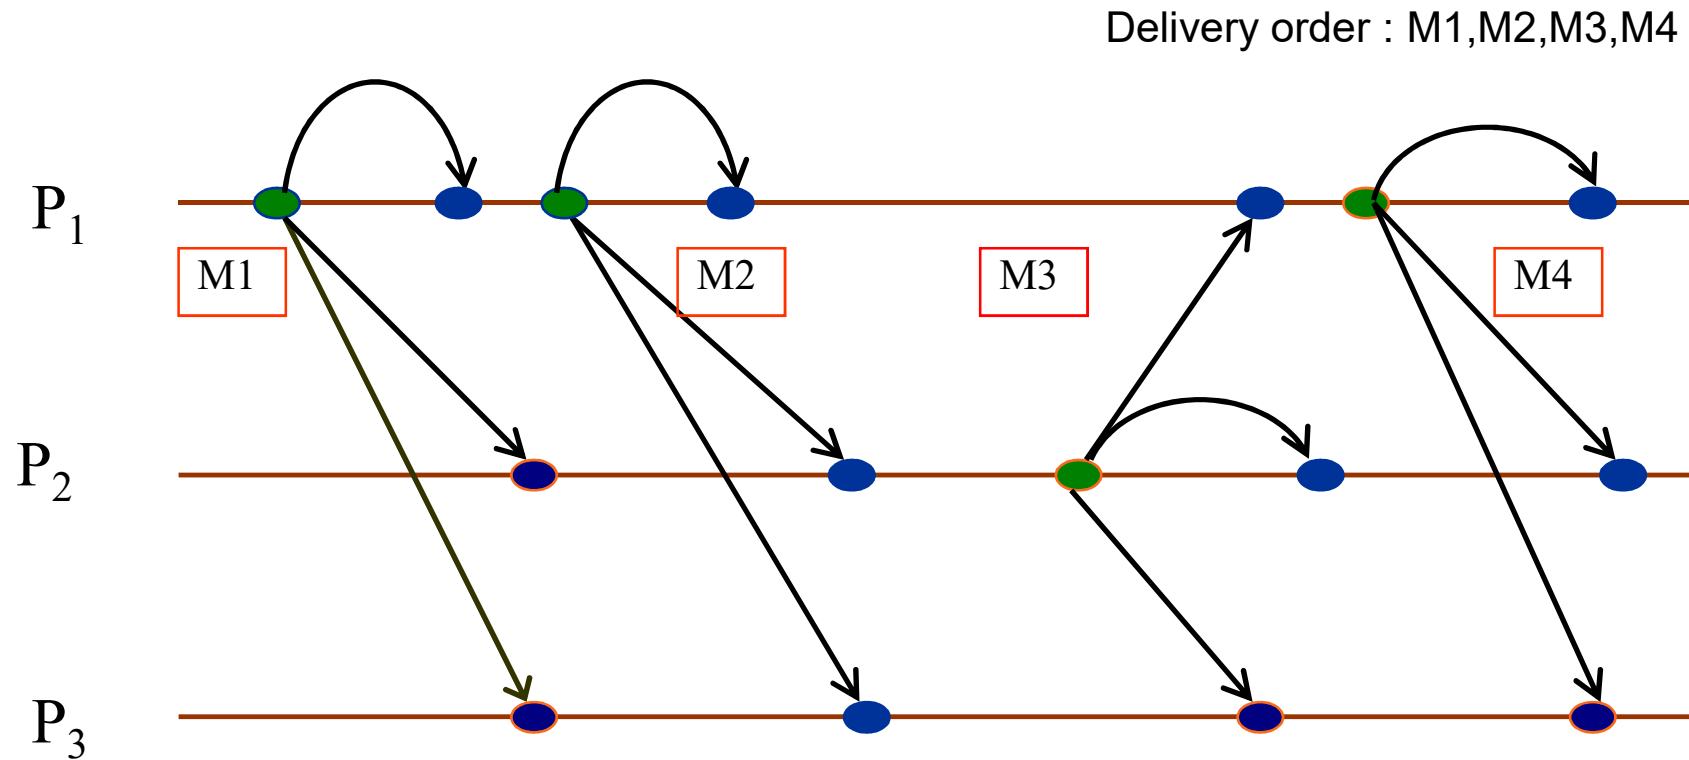
If a process delivers a message  $M_1$  before it broadcasts a message  $M_2$ , then each recipient process delivers  $M_1$  before  $M_2$ .

□ Message  $M_1$  ( ..... ) shows violation of global causal ordering.



# Causal Order

## Ordering Properties(3)



If the broadcast of a message  $M_1$  *causally precedes* the broadcast of a message  $M_2$ , then no process delivers  $M_2$  unless it has previously delivered  $M_1$ .

$M_1 \prec M_2$

$M_2 \prec M_3$

$M_1 \prec M_2 \wedge M_2 \prec M_3 \Rightarrow M_1 \prec M_3$

# SIMPLE SOLUTION TO ENSURE DELIVERY OF MESSAGES IN CAUSAL ORDER

## BIRMAN-SCHIPER-STEPHENSON PROTOCOL

### REFERENCE:

1. BIRMAN, JOSHEPH, "RELIABLE COMMUNICATION IN PRESENCE OF FAILURE", ACM TRANSACTIONS ON COMPUTER SYSTEMS, 5(1), 1987.
2. BIRMAN, SCHIPHER, STEPHENSON, "LIGHWEIGHT CAUSAL AND ATOMIC GROUP MULTICAST". ACM TRANSACTIONS ON COMPUTER SYSTEMS, 9(3), 1991

# BIRMAN-SCHIPER-STEPHENSON PROTOCOL

## BASIC IDEA

BASIC IDEA OF PROTOCOL IS TO DELIVER A MESSAGE TO A PROCESS ONLY IF THE MESSAGE IMMEDIATELY PRECEDING IT HAS BEEN DELIVERED

OTHERWISE THEY ARE BUFFERED UNTIL PRECEDING MESSAGE ARE DELIVERED.

A VECTOR ACCOMPANYING EACH MESSAGE CONTAIN INFORMATION FOR A RECIPIENT PROCESS TO DECIDE IF ALL PRECEDING MESSAGE HAS BEEN DELIVERED.

## VECTOR CLOCKS : MODIFIED UPDATE RULES

Each process  $P_i$  has a clock  $C_i$ , which is a vector of size  $n$

The clock  $C_i$  assigns a vector  $C_i(a)$  to any event  $a$  at  $P_i$

### Update rules (Clock is updated only on Message Events)

**[IR1 Broadcast Rule]** : Clock  $C_i$  is incremented when a message is **broadcast** by a process  $P_i$

$$C_i[i] = C_i[i] + d$$

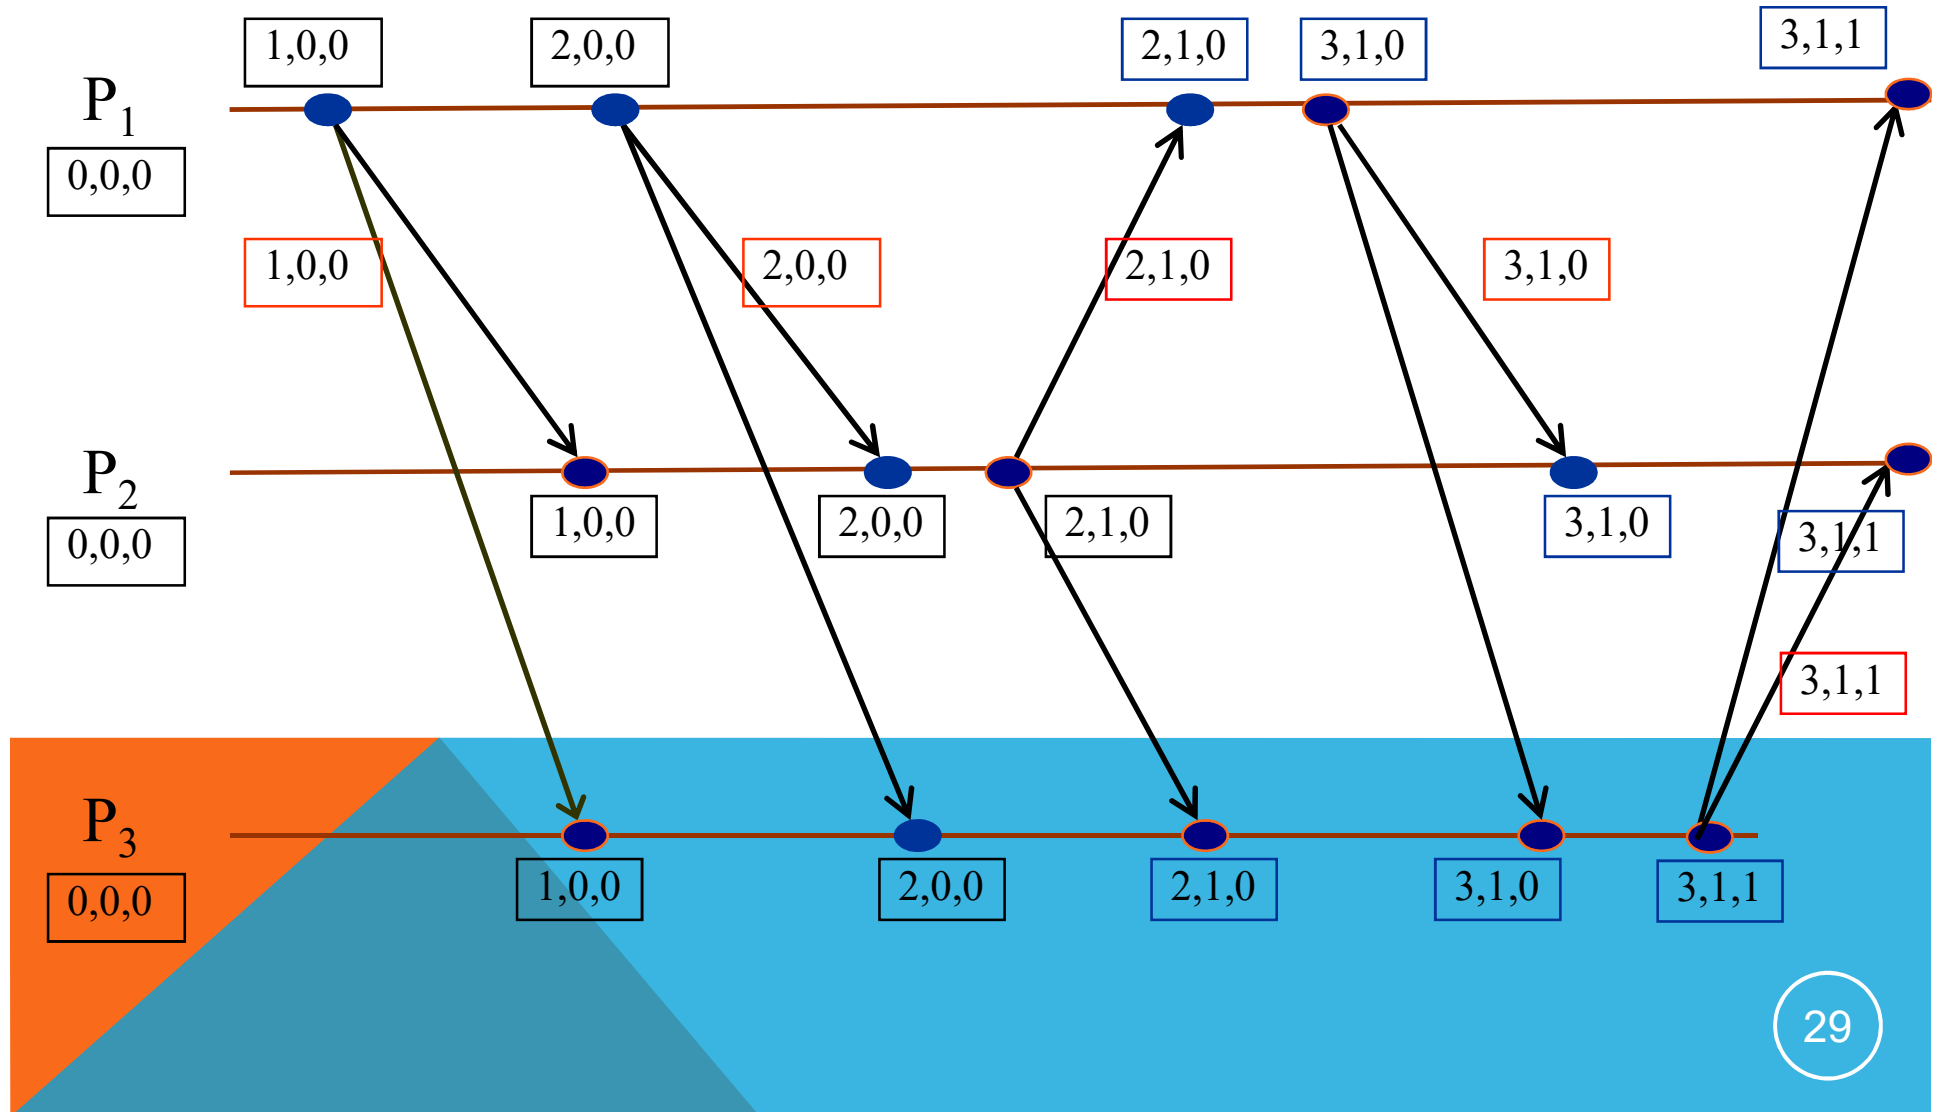
### **[IR2 Delivery Rule]**

When  $P_i$  sends a message  $m$  to Process  $P_j$ , it piggybacks a logical timestamp  $t$  which equals the time of the send event  
:  $t(m) = C_i$

When a **message is received** by a process  $P_j$  where a message with timestamp  $t$  is received, the clock is advanced.

- $C_j[k] = \max(C_j[k], t_m[k])$  for all  $k$
- **(Please note the change)**

# APPLYING VECTOR CLOCKS TO BROADCAST SYSTEM (AS PER UPDATED RULES)



# PROBLEMS OF VECTOR CLOCK

- Message size increases since each message needs to be tagged with the vector
- Size can be reduced in some cases by only sending values that have changed

## Some Good Observations

In a system where vector clock is updated on message send and receive events only

- ❑  $VT_i[i]$  indicates number of messages sent by process  $P_i$ .
- ❑  $VT_j[i]$  indicates number of messages received by process  $P_j$  sent by process  $P_i$ .


# BIRMAN-SCHIPER-STEPHENSON PROTOCOL

- To broadcast  $M$  from process  $P_i$ , increment  $VT_{Pi}[i]$  as  $VT_{Pi}[i] = VT_{Pi}[i] + 1$  and timestamp  $M$  with  $VT_m = VT_{Pi}$ .  
( Note :  $VT_{Pi}[i] - 1$  indicate how many messages from  $P_i$  precedes  $M$ )
- When  $j \neq i$  receives  $m$ ,  $j$  delays delivery of  $m$  until
  - $VT_{Pj}[i] = VT_m[i] - 1$  and
  - $VT_{Pj}[k] \geq VT_m[k]$  for all  $k \neq i$  ( i.e.,  $\forall K. K \in \{1,2,3,\dots,n\} - \{i\}$  )
  - Delayed messages are queued in  $j$  sorted by vector time. Concurrent messages are sorted by receive time.

- When  $M$  is delivered at  $P_j$ ,  $VT_{Pj}$  is updated according to vector clock rule.

$$C_j[k] = \max(C_j[k], t_m[k]) \text{ for all } k$$

# FIFO ORDER




If a message  $M_1$  is delayed and  $M_2$  arrives earlier,  $M_2$  will not be delivered as  $V_{TM2}[1]-1$  is not equal to  $V_{TP3}[1]$

□ Message  $M_1$  ( ..... ) shows violation of FIFO order.

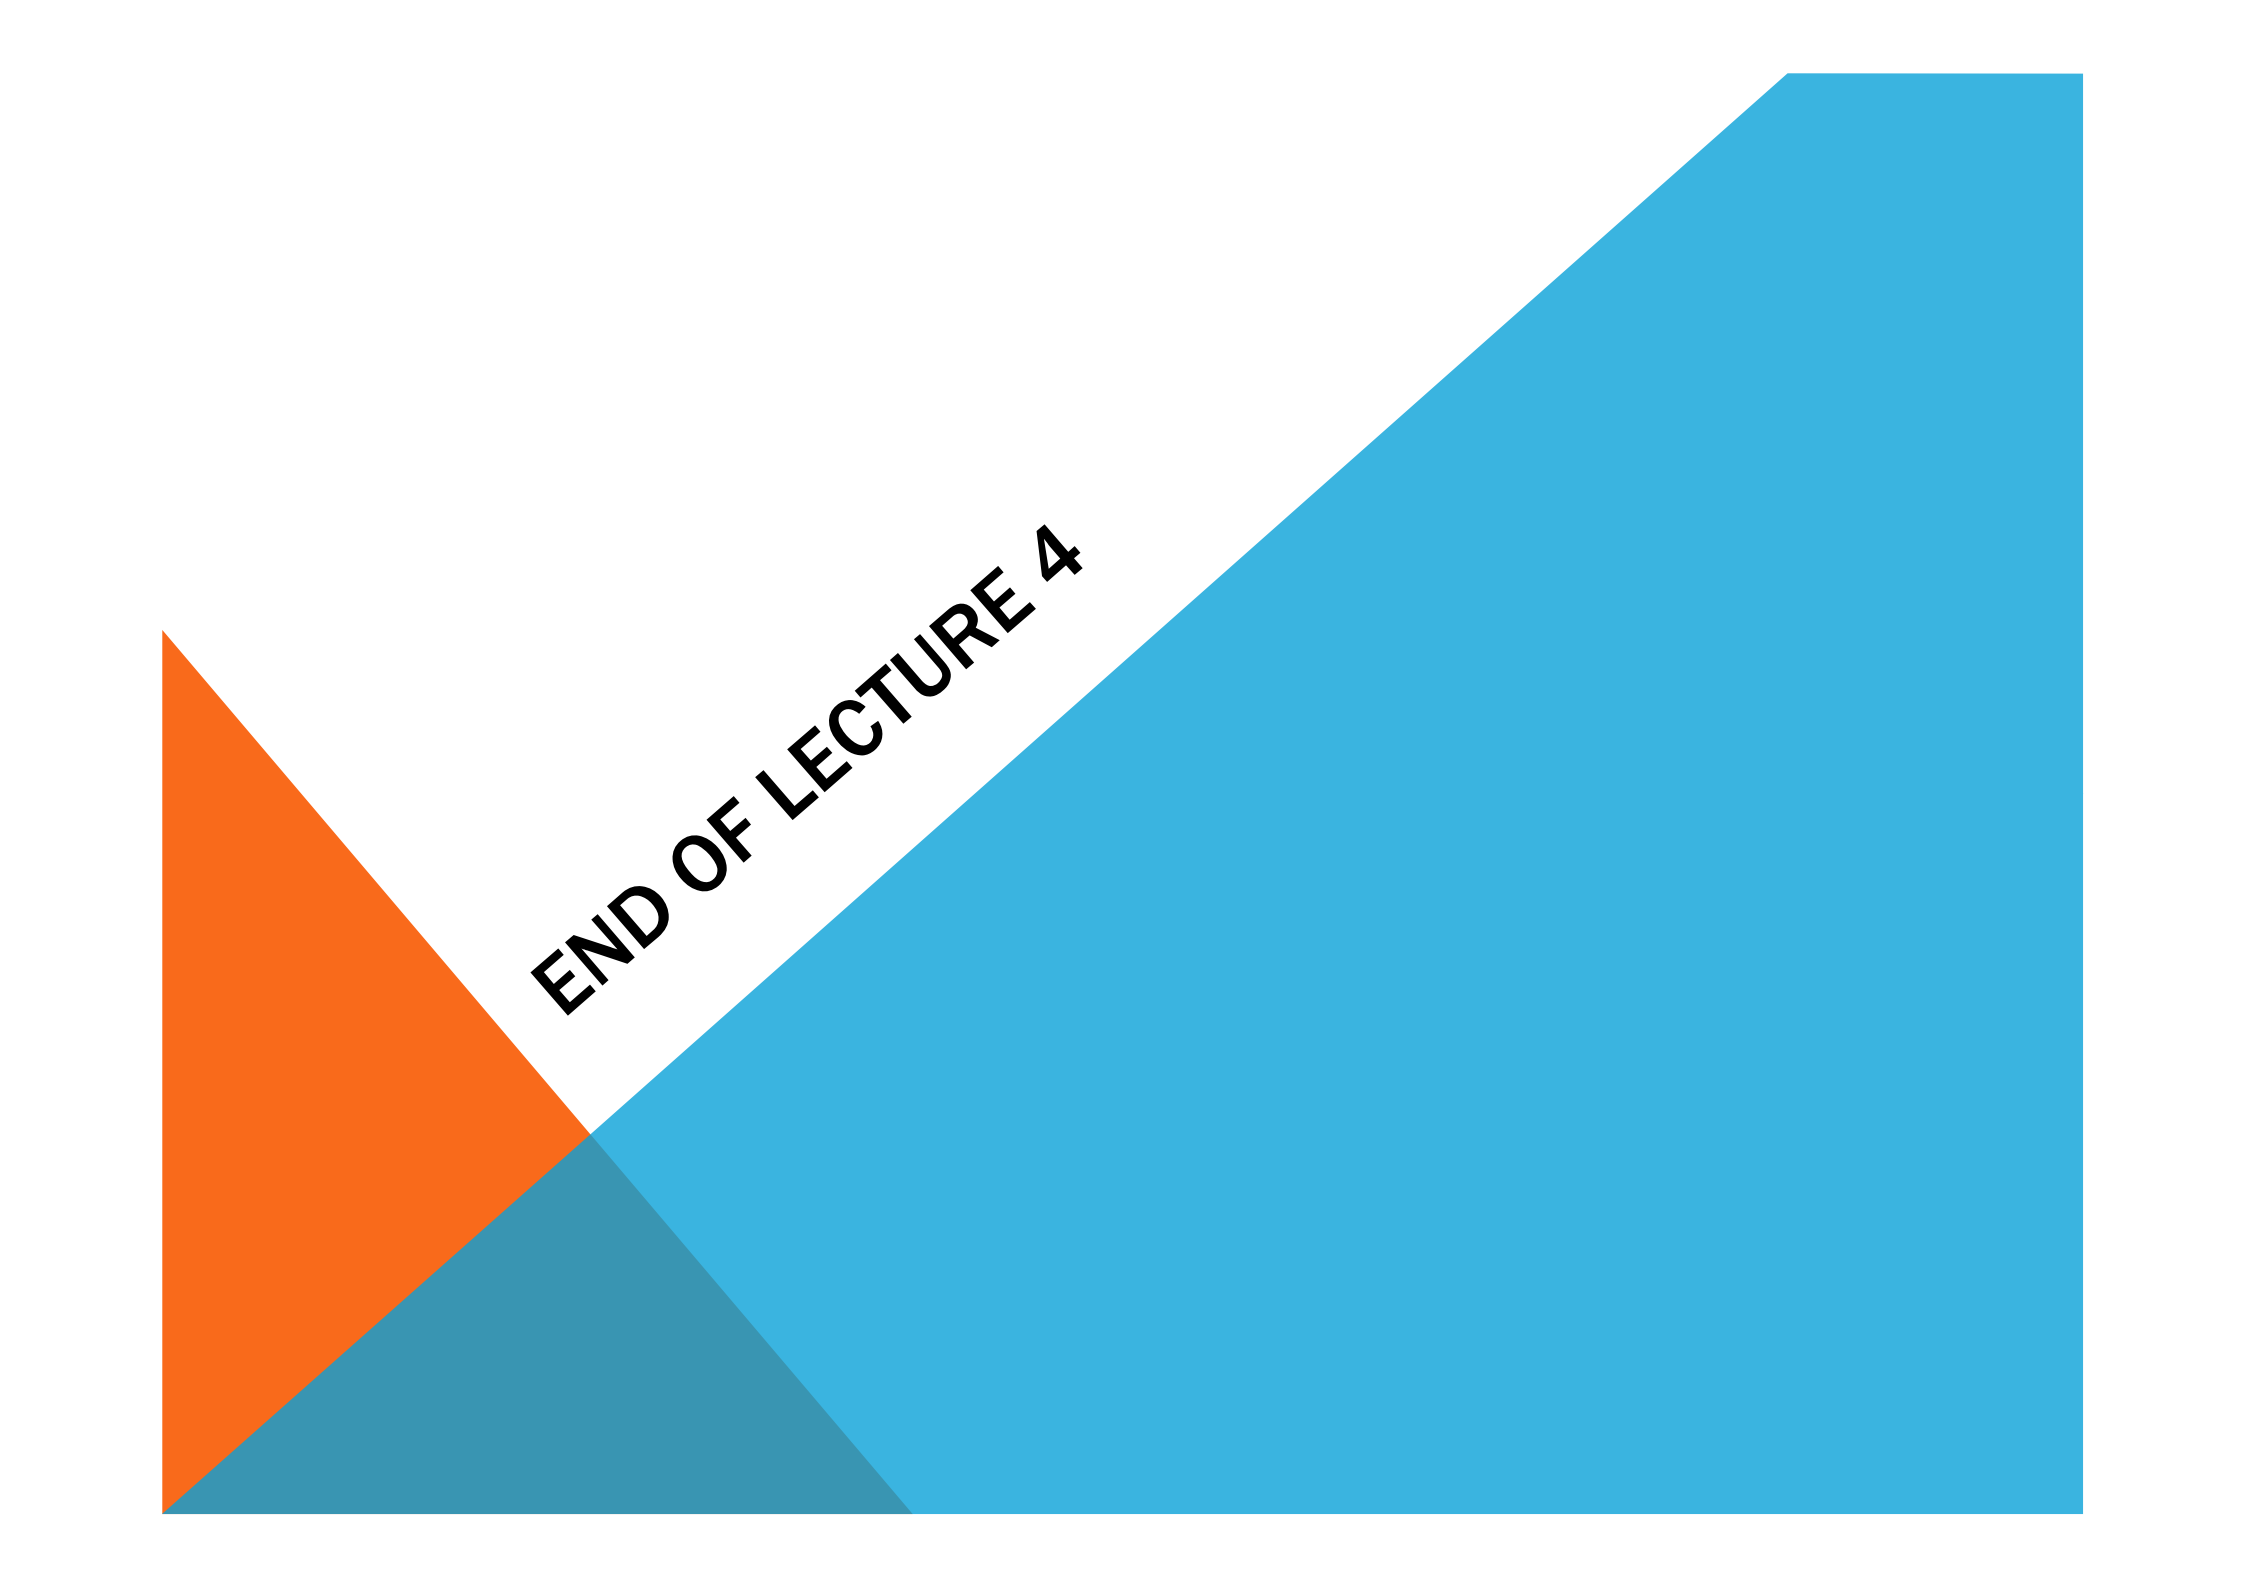


# Local Order



If a process  $P_3$  receives  $M_2$  before  $M_1$ , then  $M_2$  will be buffered until receipt of  $M_1$ . As  $VTM_2 = [1,1,0]$  &  $VTP_3 [0,0,0]$  ( $VTP_3[1] < VTM_2[1]$ ).

□ Message  $M_1$  ( ..... ) shows violation of global causal ordering.

The background features a large light blue shape on the right side. On the left, there is an orange triangle pointing towards the center, and a teal triangle pointing towards the bottom left. The text "END OF LECTURE 4" is written in black, bold, uppercase letters, rotated diagonally, and positioned in the white space between the orange and teal shapes.

**END OF LECTURE 4**